Parallel Programming using OpenMP

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University
Computer Graphics

OpenMP Multithreaded Programming

- OpenMP stands for "Open Multi-Processing"

- OpenMP is a multi-vendor (see next page) standard to perform shared-memory multithreading

- OpenMP is both compiler-directive- and library-based

- OpenMP threads share a single executable, a single global memory, and a single heap (malloc, new)

- Each OpenMP thread has its own stack (function arguments, function return address, local variables)

- Using OpenMP usually requires no dramatic code changes

- OpenMP probably gives you the biggest multithread benefit per amount of work you have to put in to using it

Much of your use of OpenMP will be accomplished by issuing C/C++ "pragmas" to tell the compiler how to build the threads into your executable, like this:

#pragma omp directive [clause]

Oregon State University
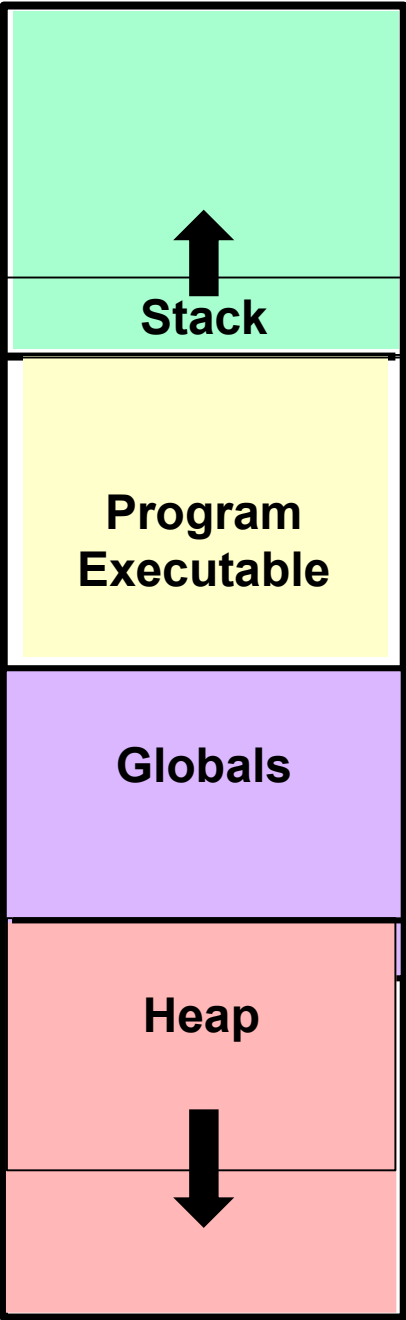Computer Graphics

Who is in the OpenMP Consortium?

What OpenMP Isn't:

- OpenMP doesn't check for data dependencies, data conflicts, deadlocks, or race conditions. You are responsible for avoiding those yourself

- OpenMP doesn't check for non-conforming code sequences (we'll talk about what this means)

- OpenMP doesn't guarantee *identical* behavior across vendors or hardware, or even between multiple runs on the same vendor's hardware

- OpenMP doesn't guarantee the *order* in which threads execute, just that they do execute

- OpenMP is not overhead-free

- OpenMP does not prevent you from writing code that triggers cache performance problems (such as in false-sharing), in fact, *it makes it really easy*
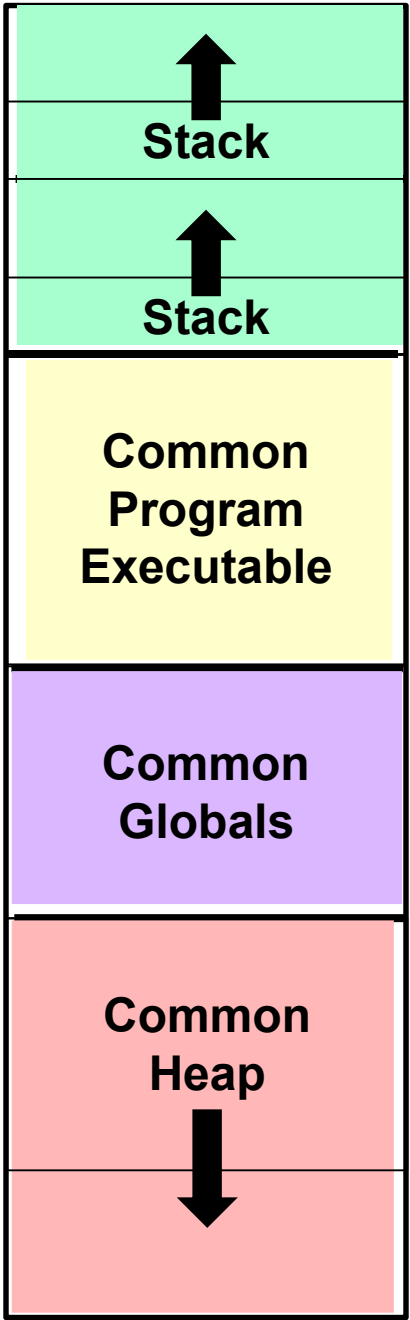
Oregon State
University
Computer Graphics

We will get to "false sharing" in the cache notes

Memory Allocation in a Multithreaded Program

One-thread

Stack
Program Executable
Globals
Heap

Don't take this completely literally. The exact arrangement depends on the operating system and the compiler. For example, sometimes the stack and heap are arranged so that they grow towards each other.

Multiple-threads

Stack
Stack
Common Program Executable
Common Globals
Common Heap

Oregon State
University
Computer Graphics

Using OpenMP on Linux

g++ -o proj proj.cpp -lm -fopenmp

Using OpenMP in Microsoft Visual Studio

1. Go to the Project menu ⟶ Project Properties

2. Change the setting Configuration Properties ⟶ C/C++ ⟶ Language ⟶ OpenMP Support to **"Yes (/openmp)"**

If you are using Visual Studio and get a compile message that looks like this:

1>c1xx: error C2338: two-phase name lookup is not supported for C++/CLI, C++/CX, or OpenMP; use /Zc:twoPhase-

then do this:

1. Go to "Project Properties"⟶ "C/C++" ⟶ "Command Line"
2. Add **/Zc:twoPhase-** in "Additional Options" in the bottom section
3. Press OK

Oregon State
University
Computer Graphics

mjb – April 9, 2024

Seeing if OpenMP is Supported on Your System

```
#ifdef   _OPENMP
    fprintf( stderr, "OpenMP release %d is supported here\n", _OPENMP );

#else
    fprintf( stderr, "OpenMP is not supported here – sorry!\n" );
    return 1;

#endif
```

Printing _OPENMP gives you a year and month of the OpenMP release that you are using

To get the OpenMP version number from the year and month, check here:

OpenMP 5.0 – November 2018

OpenMP 4.5 – November 2015

OpenMP 4.0 – July 2013

OpenMP 3.1 – July 2011

OpenMP 2.0 – March 2002

OpenMP 1.0 – October 1998

- By default, flip uses g++ 11.4, which uses OpenMP version 4.5
- Visual Studio 2022 uses OpenMP 2.0

Oregon State
University
Computer Graphics

Numbers of OpenMP threads

How to specify how many OpenMP threads you want to have available:

omp_set_num_threads(num);

Asking how many cores this program has access to:

num = omp_get_num_procs( ); ← Actually returns the number of hyperthreads,
not the number of *physical* cores

Setting the number of available threads to the exact number of cores available:

omp_set_num_threads(omp_get_num_procs( )   );

Asking how many OpenMP threads this program is using right now:

num = omp_get_num_threads( );

Asking which thread number this one is:

me = omp_get_thread_num( );

Oregon State
University
Computer Graphics

Creating an OpenMP Team of Threads

```
#pragma omp parallel default(none)
{

           . . .

}
```

This creates a
team of threads

Each thread then executes all
lines of code in this block.

Think of it this way:


```
#pragma omp parallel default(none)
```

The OpenMP Thread Team Prints a Friendly Message

```c
#include <stdio.h>
#include <omp.h>
int
main( )
{
        omp_set_num_threads( 8 );
        #pragma omp parallel default(none)
        {
                printf( "Hello, World, from thread #%d ! \n" , omp_get_thread_num(  )   );
        }
        return 0;

}
```

Hint: run it several times in a row. What do you see? Why?

Oregon State
University
Computer Graphics

Uh-oh…

First Run

Hello, World, from thread #6 !
Hello, World, from thread #1 !
Hello, World, from thread #7 !
Hello, World, from thread #5 !
Hello, World, from thread #4 !
Hello, World, from thread #3 !
Hello, World, from thread #2 !
Hello, World, from thread #0 !

Second Run

Hello, World, from thread #0 !
Hello, World, from thread #7 !
Hello, World, from thread #4 !
Hello, World, from thread #6 !
Hello, World, from thread #1 !
Hello, World, from thread #3 !
Hello, World, from thread #5 !
Hello, World, from thread #2 !

Third Run

Hello, World, from thread #2 !
Hello, World, from thread #5 !
Hello, World, from thread #0 !
Hello, World, from thread #7 !
Hello, World, from thread #1 !
Hello, World, from thread #3 !
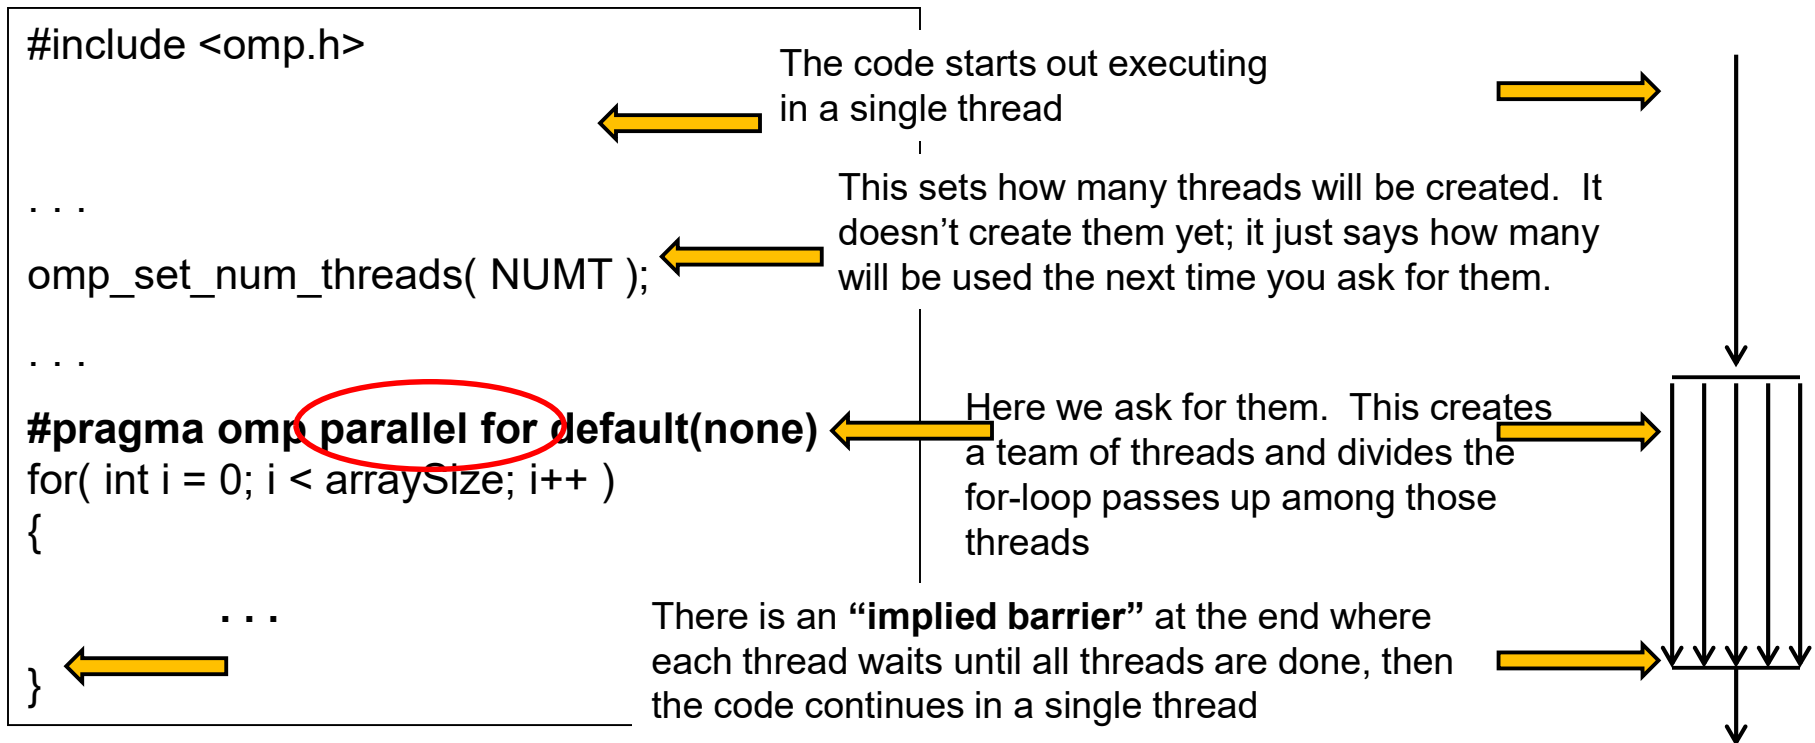Hello, World, from thread #4 !
Hello, World, from thread #6 !

Fourth Run

Hello, World, from thread #1 !
Hello, World, from thread #3 !
Hello, World, from thread #5 !
Hello, World, from thread #2 !
Hello, World, from thread #4 !
Hello, World, from thread #7 !
Hello, World, from thread #6 !
Hello, World, from thread #0 !

There is no guarantee of thread execution order!

Oregon State University
Computer Graphics

```
#include <omp.h>



. . .

omp_set_num_threads( NUMT );

. . .

#pragma omp parallel for default(none)
for( int i = 0; i < arraySize; i++ )
{

        . . .

}
```

The code starts out executing in a single thread

This sets how many threads will be created. It doesn't create them yet; it just says how many will be used the next time you ask for them.

Here we ask for them. This creates a team of threads and divides the for-loop passes up among those threads

There is an **"implied barrier"** at the end where each thread waits until all threads are done, then the code continues in a single thread

This tells the compiler to parallelize the for-loop into multiple threads. Each thread automatically gets its own personal copy of the variable *i* because it is defined within the for-loop body.

The **default(none)** directive forces you to explicitly declare all variables declared outside the parallel region to be either private or shared while they are in the parallel region. Variables declared within the for-loop are automatically private.

Computer Graphics

OpenMP for-Loop Rules

```
#pragma  omp  parallel for  default(none), shared(…), private(…)

for( int index = start ; index terminate condition; index changed )
```

- The *index* must be an *int* or a *pointer*

- The *start* and *terminate* conditions must have compatible types

- Neither the *start* nor the *terminate* conditions can be changed during the execution of the loop

- The *index* can only be modified by the *changed* expression (i.e., not modified inside the loop itself)

- You cannot use a *break* or a *goto* to get out of the loop

- There can be no inter-loop data dependencies such as:
  a[i] = a[i-1] + 1.;

a[101] = a[100] + 1.;   // what if this is the *last* line of thread #0's work?
_____

a[102] = a[101] + 1.;   // what if this is the *first* line of thread #1's work?

OpenMP For-Loop Rules

for(index = start ;

index <   end
index <= end   ;
index >   end
index >= end

index++
++index
index--
--index
index += incr
index = index + incr
index = incr + index
index -= decr
index = index - decr

)

Oregon State
University
Computer Graphics

```
float x = 0.;
#pragma omp parallel for …
for( int i = 0; i < N; i++ )
{
          x = (float) i;
          float y = x*x;
          << more code… >

}
```

i and **y** are automatically *private* because they are defined within the loop.

Good practice demands that **x** be explicitly declared to be shared or private!

private(x)
Means that each thread will get its own version of the variable

shared(x)
Means that all threads will share a common version of the variable

default(none)
I recommend that you include this in your OpenMP for-loop directive. This will force you to explicitly flag all of your externally-declared variables as *shared* or *private*. Don't make a mistake by leaving it up to the default!

Oregon State University
Computer Graphics

Example:
#pragma omp parallel for default(none), private(x)

Because of the loop dependency, this whole thing is not parallelizable:

```
x[ 0 ] = 0.;
y[ 0 ] *= 2.;
for( int i = 1; i < N; i++ )
{
        x[ i ] = x[ i-1 ] + 1.;
        y[ i ] *= 2.;
}
```

But it *can* be broken into one loop that is not parallelizable, plus one that is:

```
x[ 0 ] = 0.;
for( int i = 1; i < N; i++ )
{
        x[ i ] = x[ i-1 ] + 1.;
}

#pragma omp parallel for shared(y)
for( int i = 0; i < N; i++ )
{
        y[ i ] *= 2.;
}
```

For-loop "Collapsing"

Uh-oh, which for-loop do you put the #pragma on?

```
for( int i = 1; i < N; i++ )
{
        for( int j = 0; j < M; j++ )
        {
                . . .
        }
}
```

Ah-ha – trick question. You put it on both!

How many for-loops to collapse into one loop

```
#pragma omp parallel for collapse(2)
for( int i = 1; i < N; i++ )
{
        for( int j = 0; j < M; j++ )
        {
                . . .
        }
}
```

Oregon State
University
Computer Graphics

Single Program Multiple Data (SPMD) in OpenMP

```
#define NUM            1000000

float A[NUM], B[NUM], C[NUM];

. . .

int total = omp_get_num_threads(  );
#pragma omp parallel default(none),shared(total)

{
    int me = omp_get_thread_num(  );

    DoWork( me, total );

}
```

```
void  DoWork( int m, int t )
{
        int first = NUM * m / t;
        int last = NUM * (m+1)/t   -   1;
        for( int i = first; i <= last; i++ )
        {
                C[ i ] = A[ i ] * B[ i ];
        }
}
```

Oregon
University
Computer Graphics

OpenMP Allocation of Work to Threads

Static Threads

• All work is allocated and assigned at runtime

Dynamic Threads

• The pool is statically assigned some of the work at runtime, but not all of it
• When a thread from the pool becomes idle, it gets a new assignment
• "Round-robin assignments"

OpenMP Scheduling

schedule(static [,chunksize])
schedule(dynamic [,chunksize])
Defaults to static
chunksize defaults to 1

Oregon State
University
Computer Graphics

OpenMP Allocation of Work to Threads

```
#pragma omp parallel for default(none),schedule(static,chunksize)
for( int index = 0 ; index < 12 ; index++ )
```

Static,1

0	0,3,6,9	**chunksize = 1**
1	1,4,7,10	Each thread is assigned one iteration, then
2	2,5,8,11	the assignments start over

Static,2

0	0,1,6,7	**chunksize = 2**
1	2,3,8,9	Each thread is assigned two iterations, then
2	4,5,10,11	the assignments start over

Static,4

0	0,1,2,3	**chunksize = 4**
1	4,5,6,7	Each thread is assigned four iterations, then
2	8,9,10,11	the assignments start over

Think of dealing for-loop passes to
threads the same way as dogs deal cards
☺

Arithmetic Operations Among Threads – A Problem

```
float sum = 0.;
#pragma omp parallel for default(none), shared(sum)
 for( int i = 0; i < N; i++ )
{
      float myPartialSum = …

      sum = sum + myPartialSum;    ⟵
}
```

• There is no guarantee when each thread will execute this line

• There is not even a guarantee that each thread will finish this line before some other thread interrupts it. (Remember that each line of code usually generates multiple lines of assembly.)

• This is non-deterministic !

Assembly code:
Load sum ⟵
Add   myPartialSum
Store sum

What if the scheduler decides to switch threads right here?

Conclusion: Don't do it this way!

Here's a trapezoid integration example.
**The partial sums are added up, as shown on the previous page.**
The integration was done 30 times.
The answer is supposed to be exactly 2.
None of the 30 answers is even close.
And not only are the answers *bad*, but they are not even consistently *bad*!

0.469635	0.398893
0.517984	0.446419
0.438868	0.431204
0.437553	0.501783
0.398761	0.334996
0.506564	0.484124
0.489211	0.506362
0.584810	0.448226
0.476670	0.434737
0.530668	0.444919
0.500062	0.442432
0.672593	0.548837
0.411158	0.363092
0.408718	0.544778
0.523448	0.356299

Don't do it this way! We'll talk about how to do it correctly in the Trapezoid Integration noteset.

Mutual Exclusion Locks (Mutexes)

*omp_init_lock(omp_lock_t *);*
*omp_set_lock(omp_lock_t *);*
*omp_unset_lock(omp_lock_t *);*
*omp_test_lock(omp_lock_t *);*

Blocks if the lock is not available
Then sets it and returns when it is available

If the lock is not available, returns 0
If the lock is available, sets it and returns !0

(*omp_lock_t* is really an array of 4 unsigned chars)

Critical sections

#pragma omp critical
Restricts execution to one thread at a time

#pragma omp single
Restricts execution to a single thread ever

Barriers

#pragma omp barrier
Forces each thread to wait here until all threads arrive

(Note: there is an implied barrier after parallel for loops and OpenMP sections, unless the *nowait* clause is used)

Oregon State
University
Computer Graphics

Synchronization Example

```
omp_lock_t          Sync;
. . .
omp_init_lock( &Sync );



. . .

          Thread #0:                              Thread #1:
omp_set_lock( &Sync );                   omp_set_lock( &Sync );
<< code that needs the mutual exclusion >>   << code that needs the mutual exclusion >>
omp_unset_lock( &Sync );                 omp_unset_lock( &Sync );
```

Oregon State
University
Computer Graphics

Synchronization Example

```
omp_lock_t          Sync;
. . .
omp_init_lock( &Sync );



. . .

            Thread #0:                                Thread #1:
while(  omp_test_lock( &Sync ) == 0  )      while(  omp_test_lock( &Sync ) == 0  )
{                                           {
        DoSomeUsefulWork_0( );                      DoSomeUsefulWork_1( );
}                                           }
```

Oregon State
University
Computer Graphics

#pragma omp single

Restricts execution to a single thread ever. This is used when an operation only makes sense for one thread to do. Reading data from a file is a good example.
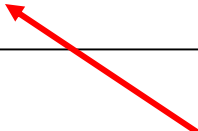
Sections are independent blocks of code, able to be assigned to separate threads if they are available.

```
#pragma omp parallel sections

{

        #pragma omp section
        {
                Task 1

        }

        #pragma omp section
        {
                Task 2

        }

}
```

(Note: there is an **implied** barrier after parallel for loops and OpenMP sections, unless the *nowait* clause is used)

A Functional Decomposition Sections Example

```
omp_set_num_threads( 3 );

#pragma omp parallel sections
{
        #pragma omp section
         {
             Watcher( );
         }

        #pragma omp section
        {
             Animals( );
        }

        #pragma omp section
        {
             Plants( );
        }

}   // implied barrier -- all functions must return to get past here
```

A Potential OpenMP/Visual Studio Compiler Problem

If you print to standard error (stderr) from inside a for-loop, like I do, then you think that you need to include *stderr* in the shared list because, well, the loops share it:

#pragma omp parallel for default(none) shared(a,b,stderr)

This turns out to be true for **g++/gcc only**.

If you are using Visual Studio, then *do not* include stderr in the list. If you do, you will get this error:

This is because:
- In g++/gcc, stderr is a *variable*
- In Visual Studio, stderr is a *defined macro*